

Tutorial 1: Threads

CS486 - Principles of Distributed Computing

Papageorgiou Spyros

Outline

- Processes
- Parallelism
- Threads overview
- Pthreads library
- Thread creation: `pthread_create`
- Waiting for threads: `pthread_join`
- Mutexes
- Barriers
- Working examples

Processes

- Carry out tasks within the operating system
- Programs
 - A set of machine code instructions and data stored on disk
- Process: a program in action
- Multiple instances of a program can exist => multiple processes
 - Each process is independent
- Have their own separate resources
 - Address space (a “chunk” of memory..)
 - Process descriptor in OS (unique identifier -> PID)

(simple_process.c)

Parallelism

- Why is it important?
- Multicore architectures prominent
 - Allow multiple contexts to run simultaneously (e.g. multiple process)
- Contexts work together in order to solve problems **faster**
- Can we use processes then?
- Yes, **but**
 - Process creation is costly (prefer long running tasks)
 - Inter-process communication (**IPC**) is slow/hard (shared memory segments, sockets, files etc)
- Ideally we need
 - A “lighter” version of a process
 - Easier communication between contexts
- **...threads do that!**

Threads overview

- A *lightweight* process that lives within a process
 - Lightweight => threads use far less resources
 - If **parent** process dies, all threads die
- Threads share the same **address space**
 - e.g. changes to global variables are visible to all threads
 - **Low communication overhead**
 - ..still must synchronize though
- Multiple threads can exist within the same process
- Threads are scheduled independently

The pthreads library

- Standardized programming interface
- Library that allows us to spawn threads
- UNIX implementation of POSIX threads => pthreads
- To use the library
 - `#include <pthread.h>`
 - Compile with `-lpthread`
- Basic API
 - `pthread_create` create a thread
 - `pthread_join` wait for a thread
 - `pthread_mutex_init` initialize a mutex
 - `pthread_mutex_lock` attempt to acquire mutex (blocking)
 - `pthread_mutex_unlock` unlock mutex
 - `pthread_mutex_destroy` destroy mutex

Thread creation: pthread_create (1 / 2)

- `int pthread_create(pthread_t *thread, const pthread_attr_t *attr, void *(*start routine) (void *), void *arg);`
- `pthread_t *thread`
 - Thread descriptor, kept by the main process to use `pthread_join` later on
- `const pthread_attr_t *attr`
 - Control thread attributes
 - Usually NULL
 - **NOTE:** can be used to control thread **affinity** (the CPU to run on)
- `void *(*start routine) (void *)`
 - Function to run when the thread is created
 - Think of it as the “main” of the thread
 - example: `void* worker_func(void *data)`

Thread creation: pthread_create (2 / 2)

- **void *arg**
 - Pass arguments to a thread
 - Can be NULL if no arguments are needed
 - Cast to void * first
 - Pass multiple arguments by creating a struct
 - **NOTE:** thread is responsible to free memory
- **Return value**
 - Returns 0 on success
 - Otherwise, non-zero value..

Waiting for threads: pthread_join

- If parent process dies (or exits) threads also exit
- Need to wait for threads to finish
- Use **pthread_join**
 - `int pthread_join(pthread_t *thread, void **retval)`
- Returns immediately if thread already exited
- If `pthread_exit` is used, threads can return meaningful values
- Return value
 - Returns 0 on success
 - Otherwise error number
- For multiple threads => loop over array holding thread descriptors passed to `pthread_create` previously

Mutexes

- Simple lock primitive
 - `pthread_mutex_t lock;`
- Initialize before use
 - `pthread_mutex_init(&lock);`
- Attempt to acquire lock. This call is blocking
 - `pthread_mutex_lock(&lock);`
- Release the lock
 - `pthread_mutex_unlock(&lock);`
- Free resources
 - `pthread_mutex_destroy(&lock);`
 - Make sure lock is not used after this call!

Barriers (1 / 2)

- Used to synchronize threads
 - e.g. make sure all threads start at the same time
- Threads are allowed to progress only when all threads have reached the same point in execution
- `pthread_barrier_t barrier;`
- Initialization
 - `int pthread_barrier_init(pthread_barrier_t *barrier,
 const pthread_barrierattr_t *restrict attr,
 unsigned count);`
 - count controls how many calls t
 - e.g. `pthread_barrier_init(&barrier, NULL, 4);`

Barriers (2 / 2)

- Waiting on a barrier
 - `pthread_barrier_wait(&barrier)`
 - Once count number of threads have called `pthread_barrier_wait`, they are allowed to progress
- Destroying a barrier
 - `pthread_barrier_destroy(&barrier)`
- See `simple_barriers.c`

Working examples

- Threads
 - `simple_threads.c`
- Barriers
 - `simple_barriers.c`
- Simple single threaded process vs multithreaded process
 - `md5_file.c`
 - `md5_file_multithreaded.c`

Thoughts on parallelization

- Amdahl's law for dummies
 - *"You are as slow as the slowest (usually sequential) part of your program"*
- Different degrees of parallelization
 - *Embarrassingly parallel*
 - Middle-ground...
 - *Inherently serial*
- Organizing data in a smart way can make things things trivially parallel
 - e.g. md5 sum calculation
 - Need lock for directory pointer (shared)
 - Multiple directories -> divide files equally -> no locking needed
- Thread placement **matters**
 - Cache pollution